


Rethinking fine-tuning to mitigate feature distortion


Aditi Raghunathan

Large-scale pretraining



Diverse (typically unlabeled) data

Step one:
pretraining



**Pretrained
model**

Step two:
adaptation



**Specialize to narrow
distribution**

Robustness to distribution shifts

A core challenge for reliable machine learning in the wild

Train



Pedestrians using a crosswalk

Deploy



Skateboarders



Important pedestrians



Pedestrians jaywalking

Distribution shifts are everywhere

Train



Deploy



Satellite remote sensing (different regions)



Train





Deploy



Wildlife conservation (different forests)




Tumor detection (new hospitals)




Sim-to-real

The generalization challenge


From scratch



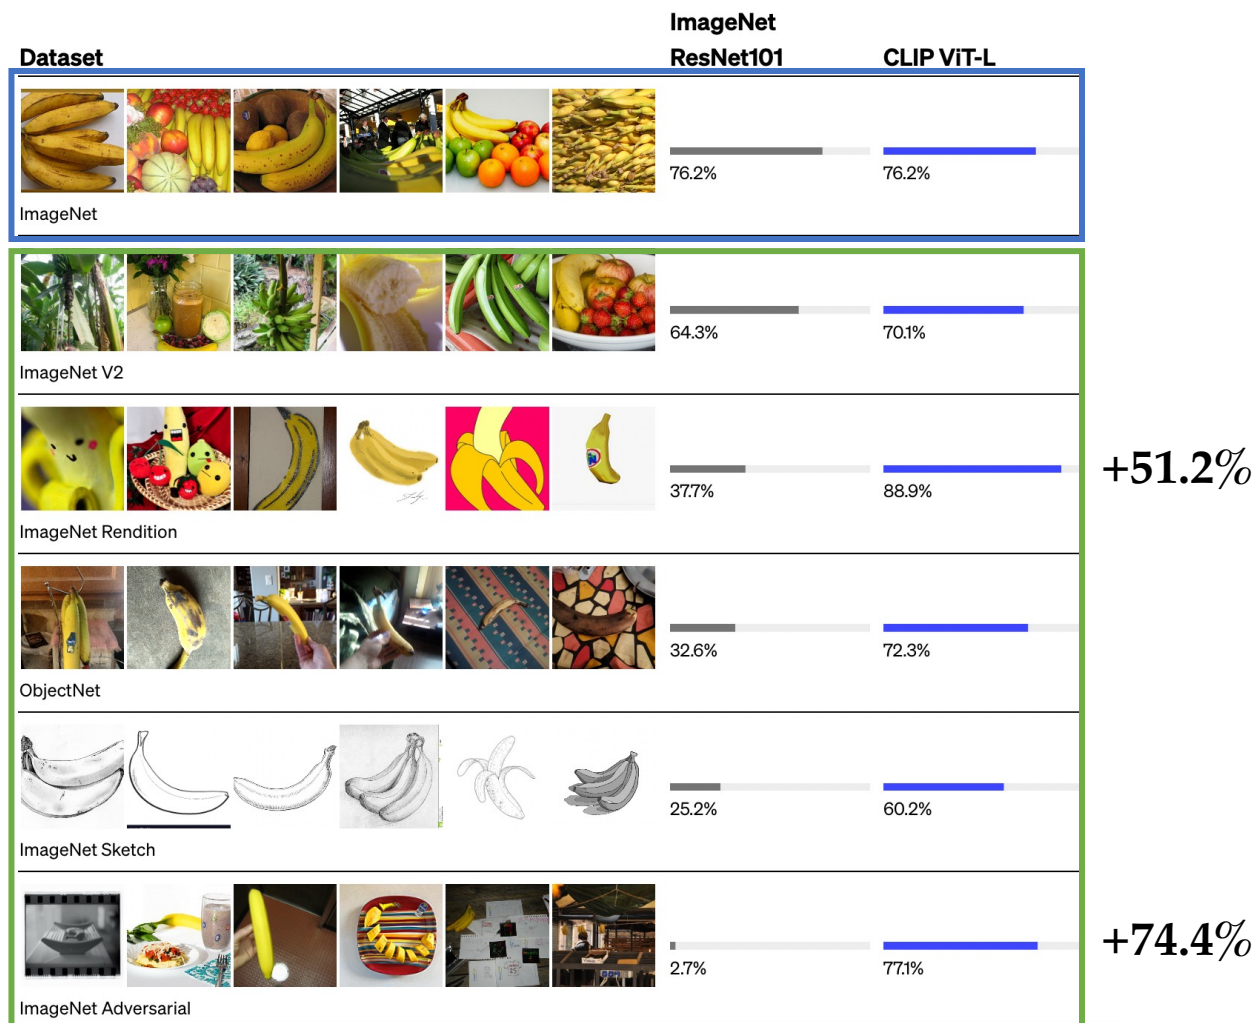
Pretraining



Fine-tuning

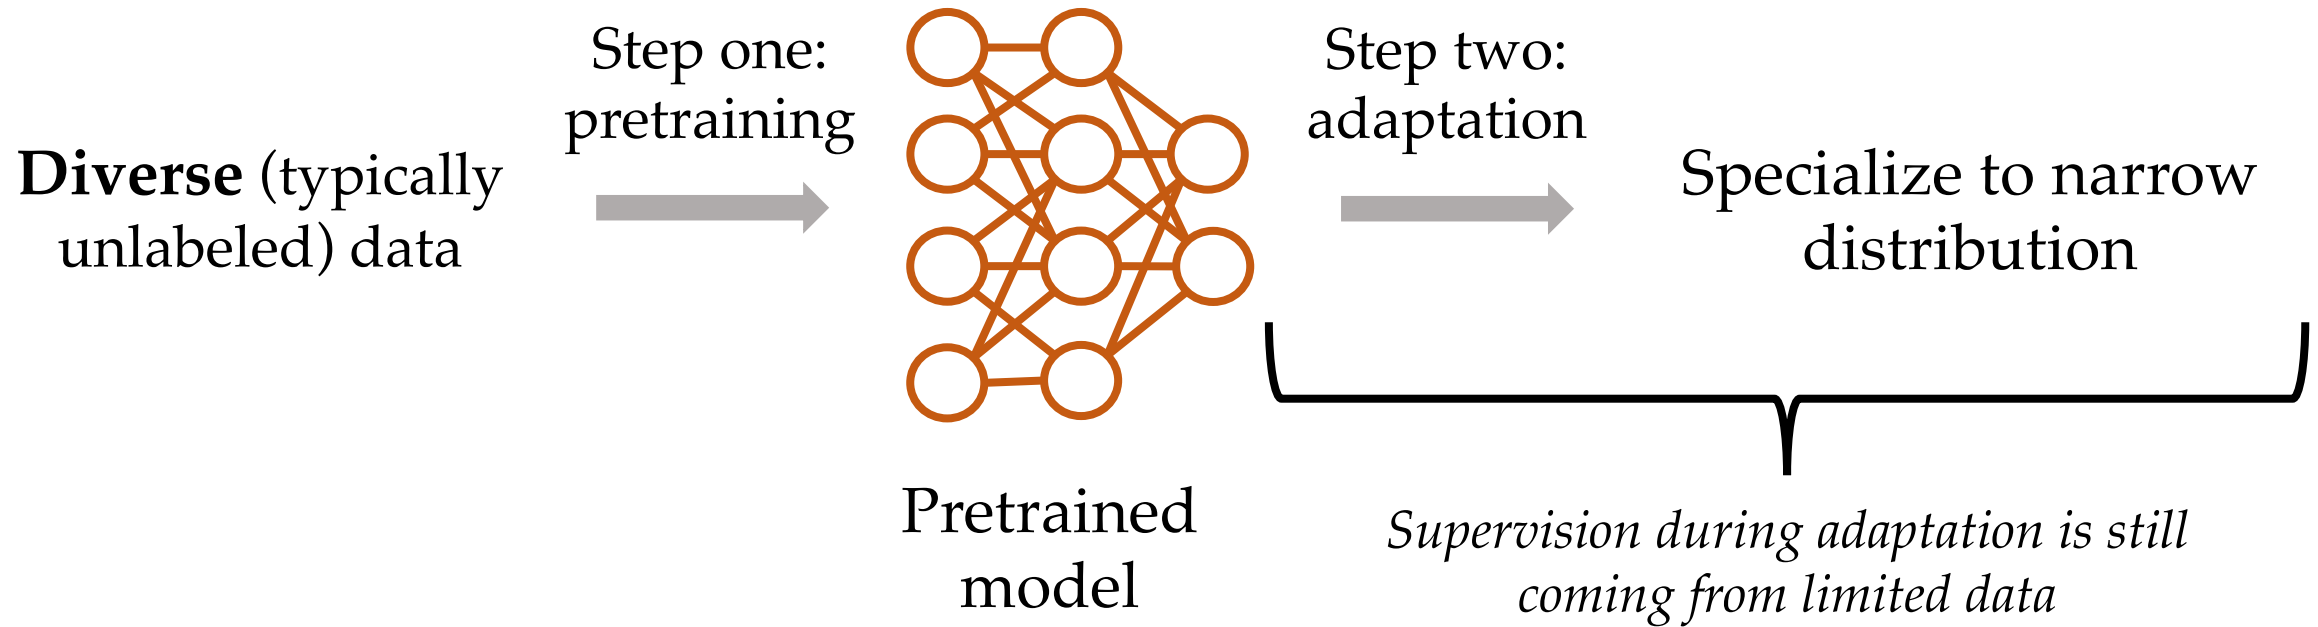


The promise of large-scale pretraining




More data
generally helps

The generalization problem revisited




The generalization challenge revisited


From scratch



Pretraining



Fine-tuning



How to retain information beyond the limited data used for adaptation?

The “art” of neural network training




- What parameters to update (model family)



- Loss function



- Optimization hyperparameters



The “art” of neural network training




- What parameters to update (model family)



- Loss function




- Optimization hyperparameters




Linear probing vs (full) fine-tuning


Pretraining



Fine-tuning



Linear probing



Pop quiz!



Dataset: BREEDS Living-17

Task: classify into animal categories

Train distribution: one subset of ImageNet hierarchy tree with animal category as root

Test distribution: other subset of ImageNet hierarchy tree with animal category as root

Pretrained model: MoCo-V2, which has seen *unlabeled* ImageNet images (including various types of animals)



Train



Test

Pop quiz: living-17

Living-17	ID	OOD
Scratch	92.4%	58.2%
Linear probing	96.5%	?
Fine-tuning	97.1%	

Does linear probing do better than scratch OOD?

Pop quiz: living-17

Living-17	ID	OOD
Scratch	92.4%	58.2%
Linear probing	96.5%	82.2%
Fine-tuning	97.1%	

Does linear probing do better
than scratch OOD?

Yes!

Pop quiz: living-17

Living-17	ID	OOD
Scratch	92.4%	58.2%
Linear probing	96.5%	82.2%
Fine-tuning	97.1%	?

Does fine-tuning do better than linear probing OOD?

Pop quiz: living-17

Living-17	ID	OOD
Scratch	92.4%	58.2%
Linear probing	96.5%	82.2%
Fine-tuning	97.1%	77.7%

Does fine-tuning do better
than linear probing OOD?

No!

Dataset: CIFAR 10.1

Task: classify into CIFAR-10 categories

Train distribution: original CIFAR-10 dataset

Test distribution: recent near-replication of the pipeline

Pretrained model: MoCo-V2, which has seen *unlabeled* ImageNet images

Pop quiz: CIFAR10.1

Living-17	ID	OOD
Linear probing	91.8%	82.7
Fine-tuning	97.3%	?

Does linear probing do better than fine-tuning OOD?

Pop quiz: CIFAR10.1


Living-17	ID	OOD
Linear probing	91.8%	82.7
Fine-tuning	97.3%	92.3%

Does linear probing do better
than fine-tuning OOD?


No!

Linear probing vs fine-tuning summary


Pretraining



Fine-tuning



Linear probing



Which method does better?

Linear probing vs fine-tuning summary

	ID	OOD
Linear probing	82.9%	
Fine-tuning	85.1%	

Averaged over 10 datasets

Common wisdom is fine-tuning works better than linear probing

Linear probing vs fine-tuning summary



	ID	OOD
Linear probing	82.9%	66.2%
Fine-tuning	85.1%	59.3%

Averaged over 10 datasets

LP performs better than FT OOD on 8 out of 10 datasets


Intuition for theoretical result

Pretrained
Features




Intuition for theoretical result

Pretrained
Features




Fine-tuning: features for ID examples change in sync with the linear head




Features for OOD examples
change less

Intuition for theoretical result

Pretrained
Features




Fine-tuning: features for ID examples change in sync with the linear head




Features for OOD examples
change less

Intuition for theoretical result

Pretrained
Features



Fine-tuning: features for ID examples change in sync with the linear head




■ ID
■ OOD

Features for OOD examples
change less




Intuition for theoretical result

Pretrained
Features



Fine-tuning: features for ID examples change in sync with the linear head




■ ID
■ OOD


Features for OOD examples
change less

Intuition for theoretical result

Pretrained
Features




Fine-tuning: features for ID
examples change in sync
with the linear head



Head performs
poorly on OOD
examples


Linear probing: freezes
pretrained features



Head is decent on
OOD examples

Key takeaway

A larger change in parameters can **distort** pretrained features



How to retain information beyond the limited data used for adaptation?

Best of both worlds

Why does FT do better ID?

Training data may not be linearly separable in the space of pre-trained features i.e. imperfect pre-trained features

Why does FT do worse OOD?


Features can change a lot to accommodate a randomly initialized head

Can we refine features without distorting them too much?


Method to achieve best of both worlds

Idea: modify pre-trained features **only as necessary**

Step 1: Linear probe



Step 2: Fine-tune



Method to achieve best of both worlds

Idea: modify pre-trained features **only as necessary**

Step 1: Linear probe

Step 2: Fine-tune

LP-FT method

Can prove that LP-FT dominates both LP and FT under the simple setting of perfect features

Improving fine-tuning

	ID	OOD	
Linear probing	82.9%	66.2%	
Fine-tuning	85.1%	59.3%	+10% over fine-tuning!
LP-FT	85.7%	68.9%	

LP-FT obtains better than the best of both worlds

The “art” of neural network training




- What parameters to update (model family)



- **Loss function**

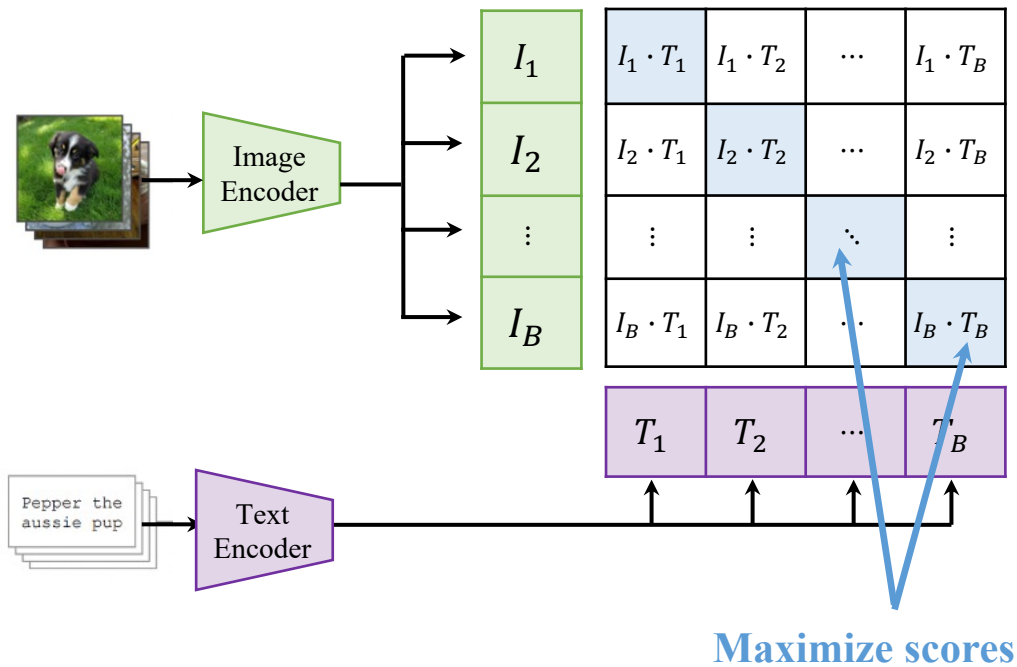


- Optimization hyperparameters




The loss function

Contrastive pretraining



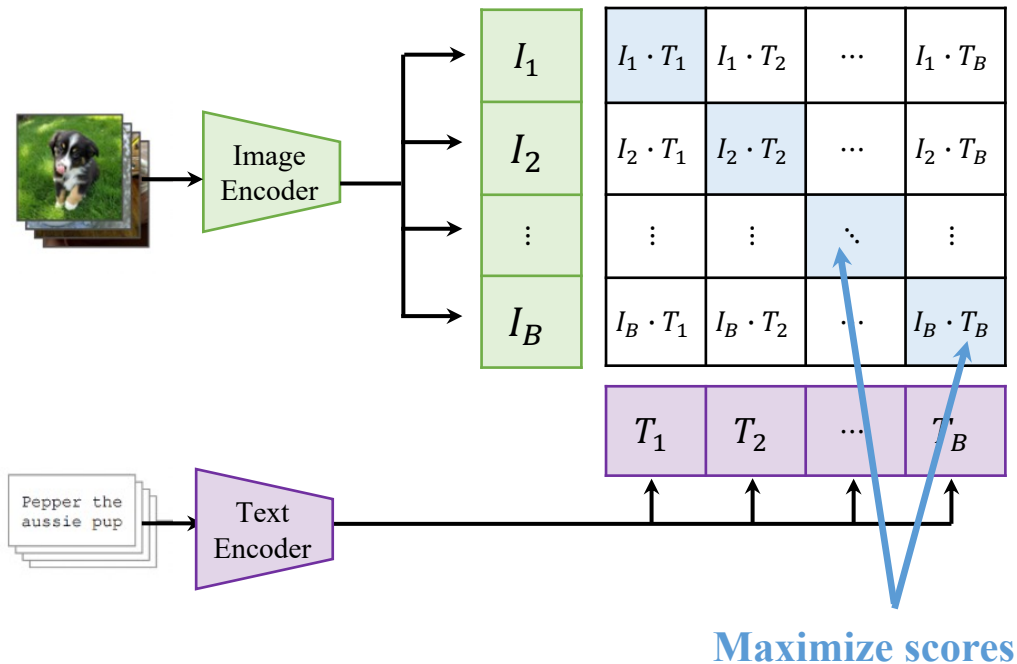
Standard finetuning



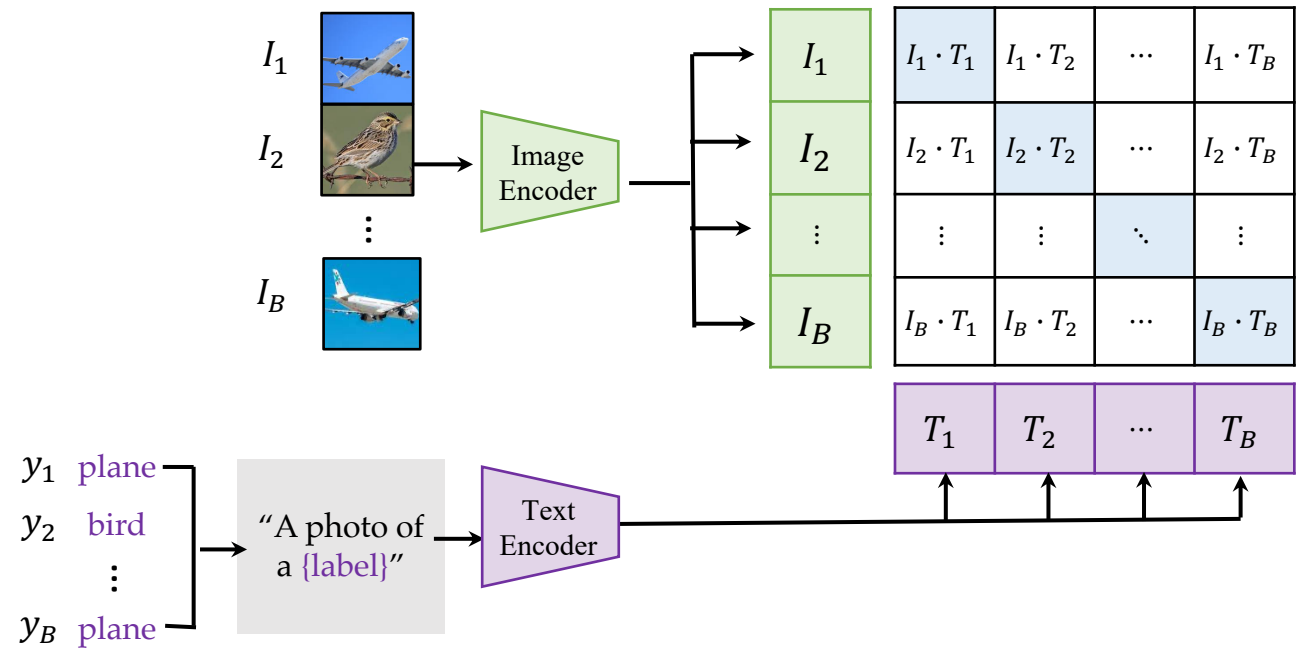
Can we reduce distortion?

Revisiting the fine-tuning loss function


Contrastive pretraining



Finetune like you pretrain (FLYP)



Fine-tune like you pretrain



Same pretraining loss can reduce distortion and improve robustness

- Full finetuning
- LP-FT
- L2-sp (baseline)
- FLYP (ours)

Fine-tune like you pretrain

Also see gains in few-shot learning

	PatchCamelyon	SST2
Zero shot	56.5%	60.5%
FT	63.1%	61.1%
LP-FT	62.7%	60.9%
FLYP	66.9%	61.3%

Summary

- Pretrained models give large improvements in accuracy, but how we fine-tune them is key
- **General principle: minimize distortion while fine-tuning**
- Two simple ways to do that
 - LP-FT (only change features once the head is trained)
 - FLYP (keep the fine-tuning loss identical to the pretraining loss)


Thanks!



Ananya Kumar



Robbie Jones



Tengyu Ma



Percy Liang

Apple

Google

Schmidt Futures

Open Philanthropy



Sachin Goyal



Sankalp Garg



Zico Kolter